

## Assignment 2 Part 2: Algorithm Description and Results

# Implementing Naive Bayes Classifier using Spark MapReduce

---

### Naïve Bayes Algorithm description:

```
function TRAIN NAIVE BAYES(D, C) returns  $\log P(c)$  and  $\log P(w|c)$ 

for each class  $c \in C$            # Calculate  $P(c)$  terms
     $N_{doc}$  = number of documents in D
     $N_c$  = number of documents from D in class c
     $\logprior[c] \leftarrow \log \frac{N_c}{N_{doc}}$ 
     $V \leftarrow$  vocabulary of D
     $bigdoc[c] \leftarrow \text{append}(d)$  for  $d \in D$  with class c
    for each word  $w$  in  $V$            # Calculate  $P(w|c)$  terms
         $count(w, c) \leftarrow$  # of occurrences of  $w$  in  $bigdoc[c]$ 
         $\loglikelihood[w, c] \leftarrow \log \frac{count(w, c) + 1}{\sum_{w' \text{ in } V} (count(w', c) + 1)}$ 
    return  $\logprior, \loglikelihood, V$ 

function TEST NAIVE BAYES( $testdoc, \logprior, \loglikelihood, C, V$ ) returns best  $c$ 

for each class  $c \in C$ 
     $sum[c] \leftarrow \logprior[c]$ 
    for each position  $i$  in  $testdoc$ 
         $word \leftarrow testdoc[i]$ 
        if  $word \in V$ 
             $sum[c] \leftarrow sum[c] + \loglikelihood[word, c]$ 
    return  $\text{argmax}_c sum[c]$ 
```

**Figure 4.2** The naive Bayes algorithm, using add-1 smoothing. To use add- $\alpha$  smoothing instead, change the +1 to + $\alpha$  for loglikelihood counts in training.

### Map Reduce Algorithm Description:

1. Log prior calculation:  
For each label, the calculate the prior probability by dividing the count of the label by the total count of all labels.  
The prior probability is added as a new column (prior) and the log prior is added as a new column (log\_prior).
2. Log Likelihood calculation:
  - Count Smoothing: The count of each word is incremented using the incr function, which likely adds a smoothing value (like Laplace smoothing) to avoid zero probabilities. The result is stored in a new column count\_smoothed.

- Total Word Count Smoothing: Similarly, the total word count for each sentiment is incremented using the `incr_v` function, and the result is stored in a new column `total_word_count_smoothed`.
- The likelihood of each word given its sentiment is calculated by dividing the smoothed word count (`count_smoothed`) by the smoothed total word count (`total_word_count_smoothed`). This is stored in a new column `likelihood`.
- The logarithm of the likelihood is computed to facilitate calculations in logarithmic space, improving numerical stability and efficiency. This is stored in a new column `log_likelihood`.

### 3. Model and result creation:

- The `test_df` DataFrame is exploded so that each word in the `words_stemmed` column gets its own row. This helps in calculating the probabilities for each word separately.
- The exploded DataFrame is joined with the `priors` DataFrame on the `sentiment` column to add the log prior probabilities.
- It is then joined with the `likelihood` DataFrame on both the `sentiment` and `word` columns to add the log likelihood probabilities for each word.
- The data is grouped by `words_stemmed`, `sentiment`, and `log_prior`.
- The sum of the `log_likelihood` values for all words in each review is calculated and stored in a new column `sum_log_likelihood`.
- The log-probability for each review is computed by adding the `log_prior` to the `sum_log_likelihood`. This value is stored in a new column `log_probability`.
- intermediate columns `log_prior` and `sum_log_likelihood` is dropped for clarity.
- The resulting DataFrame, which contains `words_stemmed` and their corresponding `log_probability`, is joined back with the original `test_df` (with `sentiment` renamed to `label`) to retain the original structure of the DataFrame.

### Priors for each class:

	$A_C^B$ sentiment	1.2 prior	1.2 log_prior
1	positive	0.5023891993022022	-0.6883801622634533
2	negative	0.4976108006977979	-0.6979370322103389

**Evaluation:** The accuracy is low as the implementation of Naïve Bayes classifier is very simple without any optimizations. For comparison, SKlearn's vanilla NB gives accuracy of 73% with TFIDF vectorization.

Accuracy	0.66
Precision	0.59
Recall	0.68
F1 measure:	0.64